

The Seductions of Scala - Applications

Dean Wampler
dean@concurrentthought.com
[@deanwampler](https://twitter.com/deanwampler)
polyglotprogramming.com/talks

April 26, 2013



I

Copyright © 2009-2013, Dean Wampler, All Rights Reserved

Monday, April 22, 13

The online version contains more material. You can also find this talk and the code used for many of the examples at [github.com/deanwampler/Presentations/tree/master/](https://github.com/deanwampler/Presentations/tree/master/SeductionsOfScala)


[SeductionsOfScala](https://github.com/deanwampler/Presentations/tree/master/SeductionsOfScala)

Copyright © 2010–2013, Dean Wampler. Some Rights Reserved – All use of the photographs and image backgrounds are by written permission only. The content is free to reuse, but attribution is requested.

<http://creativecommons.org/licenses/by-nc-sa/2.0/legalcode>

Let's examine two
popular Scala
toolkits used to
implement real-
world applications.


- Akka: Actors for highly-concurrent applications.
- Scalding: Big data the functional way.



Akka: Actor Concurrency

When you share mutable state...

Hic sunt dracones
(Here be dragons)



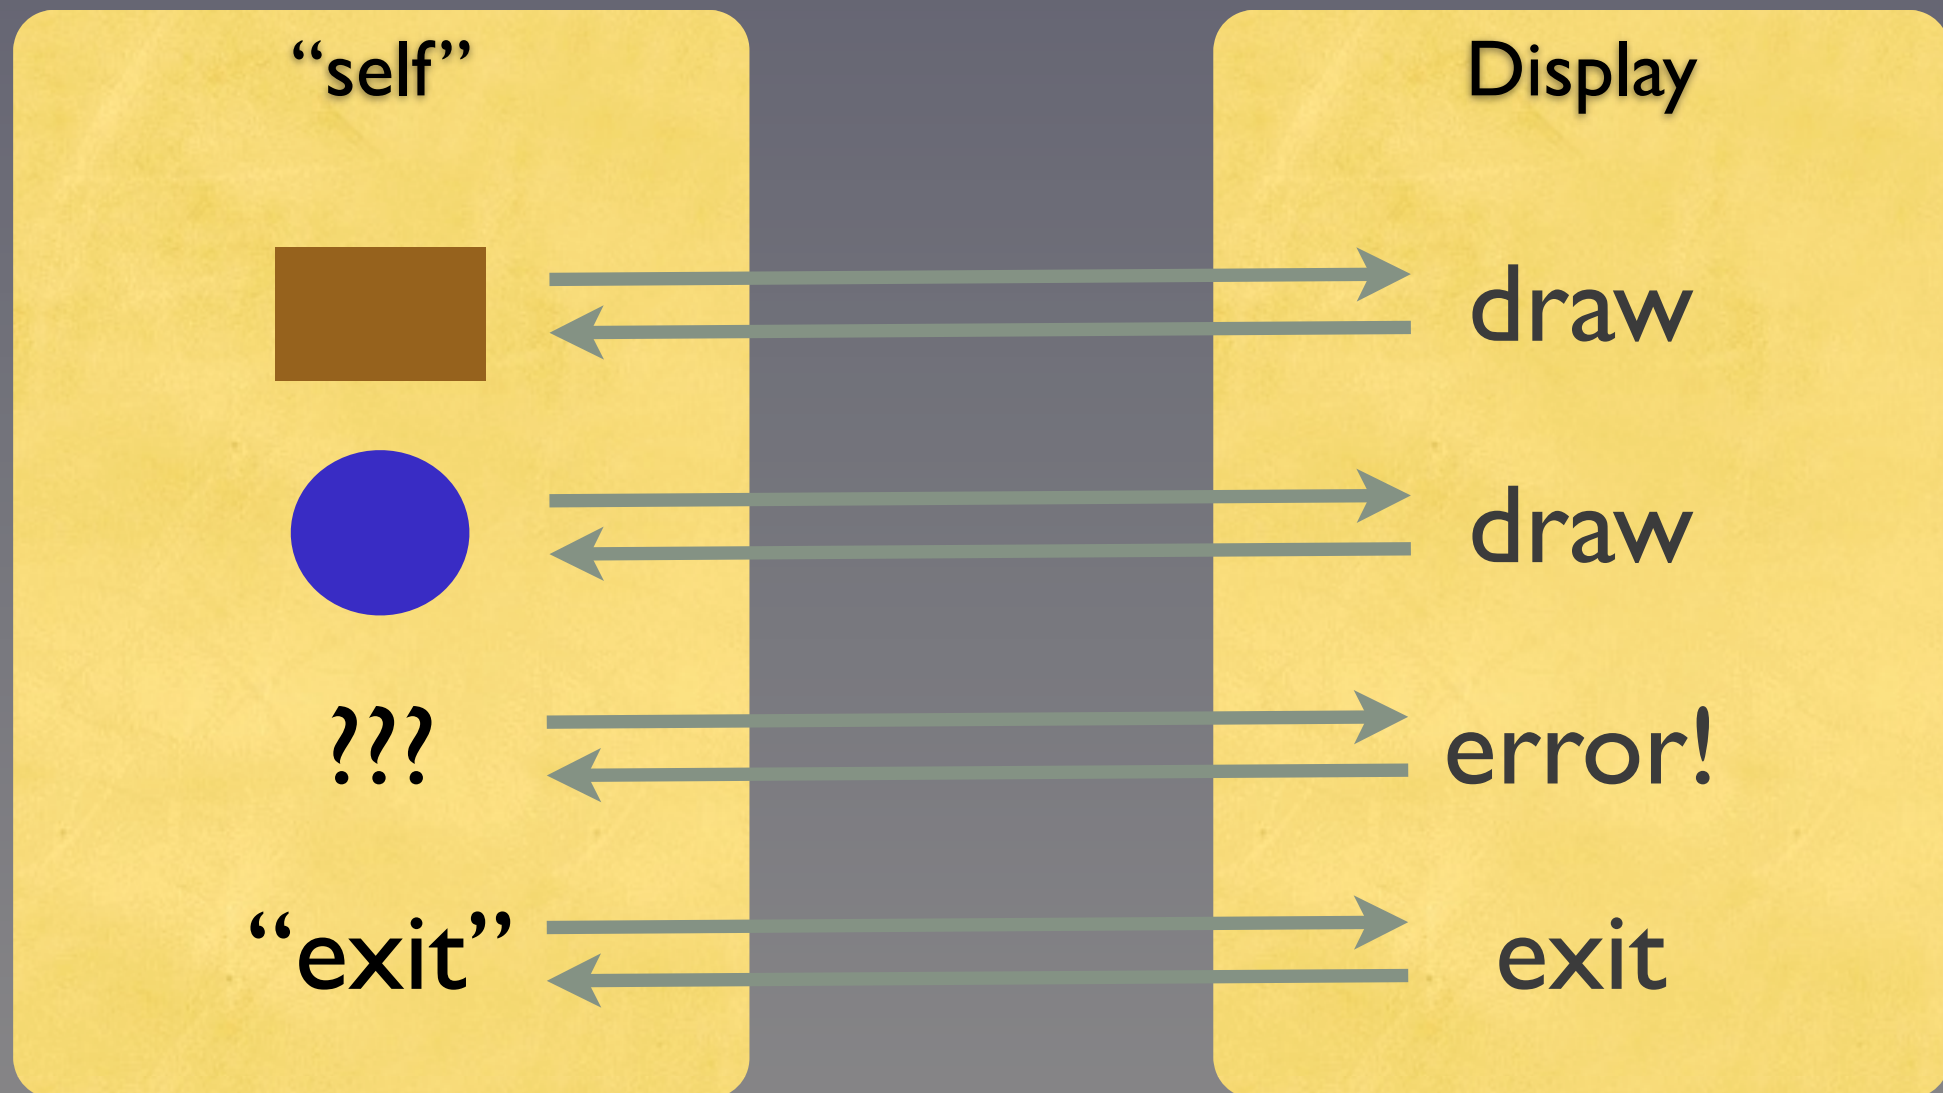
Actor Model

- Message passing between autonomous actors.
- No shared (mutable) state.

Actor Model

- First developed in the 70's by Hewitt, Agha, Hoare, etc.
- Made “famous” by Erlang.
 - Scala's Actors patterned after Erlang's.

2 Actors:




```
package shapes
```

```
case class Point(  
  x: Double, y: Double)
```

```
abstract class Shape {  
  def draw()  
}
```

*abstract **draw** method*

Hierarchy of geometric shapes

```
case class Circle(  
  center:Point, radius:Double)  
  extends Shape {  
  def draw() = ...  
}
```

*concrete **draw**
methods*

```
case class Rectangle(  
  ll:Point, h:Double, w:Double)  
  extends Shape {  
  def draw() = ...  
}
```

*Use the Akka
Actor library*

```
package shapes  
import akka.actor.Actor
```

Actor

```
class Drawer extends Actor {  
  def receive = {  
    ...  
  }  
}
```

*receive and handle
each message*

Actor for drawing shapes

||

Copyright © 2009-2013, Dean Wampler, All Rights Reserved

Monday, April 22, 13

An actor that waits for messages containing shapes to draw. Imagine this is the window manager on your computer. It loops indefinitely, blocking until a new message is received...

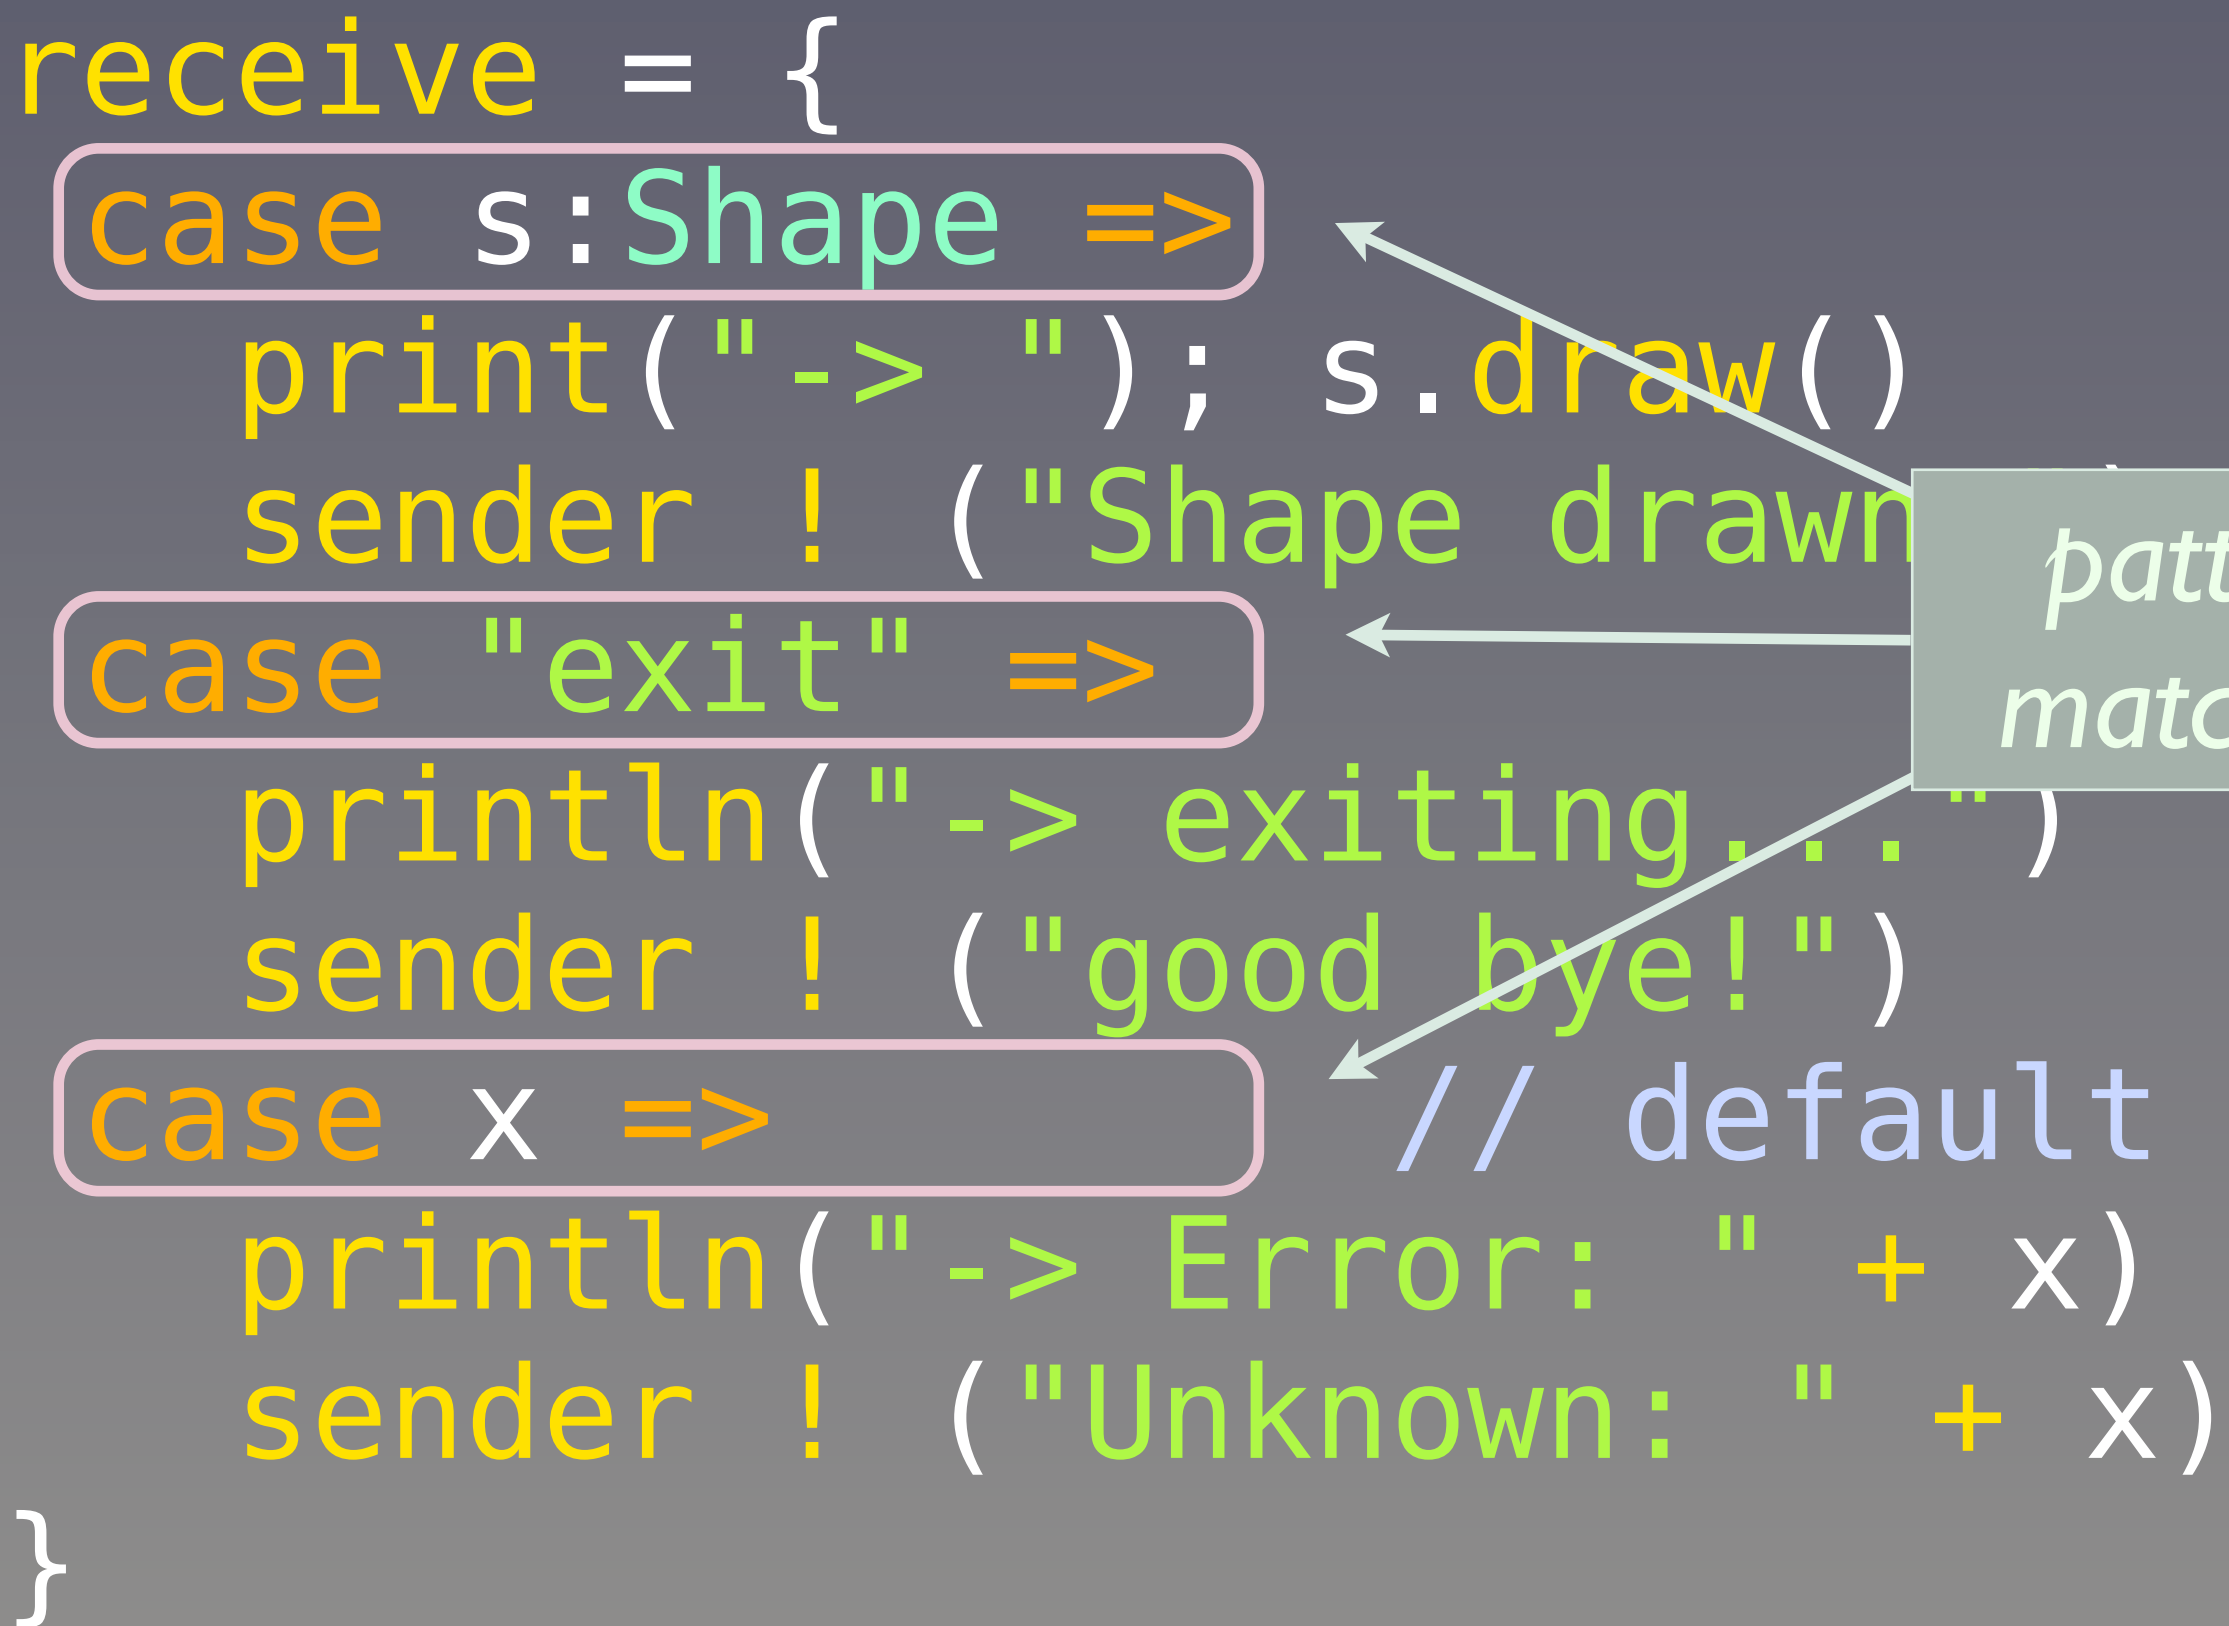
Note: This example uses the Akka Frameworks Actor library (see <http://akka.io>), which has now replaced Scala's original actors library. So, some of the basic actor classes are part of Scala's library, but we'll use the full Akka distribution.

*Receive
method*

```
receive = {  
  case s:Shape =>  
    print("-> "); s.draw()  
    sender ! ("Shape drawn.")  
  case "exit" =>  
    println("-> exiting...")  
    sender ! ("good bye!")  
  case x => // default  
    println("-> Error: " + x)  
    sender ! ("Unknown: " + x)  
}
```



```
receive = {  
  case s:Shape =>  
    print("-> "); s.draw()  
    sender ! ("Shape drawn")  
  case "exit" =>  
    println("-> exiting...")  
    sender ! ("good bye!")  
  case x => // default  
    println("-> Error: " + x)  
    sender ! ("Unknown: " + x)  
}
```



The diagram illustrates pattern matching in Scala. A box labeled "pattern matching" has three arrows pointing to the three "case" statements in the code. The first arrow points to "case s:Shape =>", the second to "case \"exit\" =>", and the third to "case x =>".

```
receive = {  
  case s:Shape =>
```

*draw shape
& send reply*

```
    print("-> "); s.draw()  
    sender ! ("Shape drawn.")
```

```
  case "exit" =>
```

```
    println("-> exiting...")  
    sender ! ("good bye!")
```

done

```
  case x => // default
```

```
    println("-> Error: " + x)  
    sender ! ("Unknown: " + x)
```

```
}
```

sender ! sends a reply

unrecognized message

```

package shapes
import akka.actor.Actor
class Drawer extends Actor {
  receive = {
    case s:Shape =>
      print("-> "); s.draw()
      sender ! ("Shape drawn.")
    case "exit" =>
      println("-> exiting...")
      sender ! ("good bye!")
    case x => // default
      println("-> Error: " + x)
      sender ! ("Unknown: " + x)
  }
}

```

Altogether

```

import shapes._
import akka.actor._
import com.typesafe.config._

object Driver {
  def main(args:Array[String])={
    val sys = ActorSystem(...)
    val driver=sys.actorOf[Driver]
    val drawer=sys.actorOf[Drawer]
    driver ! Start(drawer)
  }
}

```

Application driver

...


```
import shapes._
import akka.actor._
import com.typesafe.config._
```

```
object Driver {  
  def main(args:Array[String])={  
    val sys = ActorSystem(...)  
    val driver=sys.actorOf[Driver]  
    val drawer=sys.actorOf[Drawer]  
    driver ! Start(drawer)  
  }  
}
```

*Singleton for **main***

*Instantiate
actors*

*Send a message to
start the actors*

...

Companion class



```
...  
class Driver extends Actor {  
  val drawer: Option[Drawer] =  
    None  
  
  def receive = {  
    ...  
  }  
}
```

```
def receive = {  
  case Start(d) =>  
    drawer = Some(d)  
    d ! Circle(Point(...), ...)  
    d ! Rectangle(...)  
    d ! 3.14159  
    d ! "exit"  
  case "good bye!" =>  
    println("<- cleaning up...")  
    context.system.shutdown()  
  case other =>  
    println("<- " + other)  
}
```

sent by driver

sent by drawer

```
d ! Circle(Point(...),...)
d ! Rectangle(...)
d ! 3.14159
d ! "exit"
```

```
-> drawing: Circle(Point(0.0,0.0),1.0)
-> drawing: Rectangle(Point(0.0,0.0),
2.0,5.0)
-> Error: 3.14159
-> exiting...
<- Shape drawn.
<- Shape drawn.
<- Unknown: 3.14159
<- cleaning up...
```

*“<-” and “->” messages
may be interleaved!!*


```
...  
// Drawing.receive  
receive = {  
  case s:Shape =>  
    s.draw()  
  self.reply("...")  
  
  case ...  
  case ...  
}
```

*Functional-style
pattern matching*

*Object-
oriented-style
polymorphism*

*“Switch” statements are
not (necessarily) evil*

Exercise

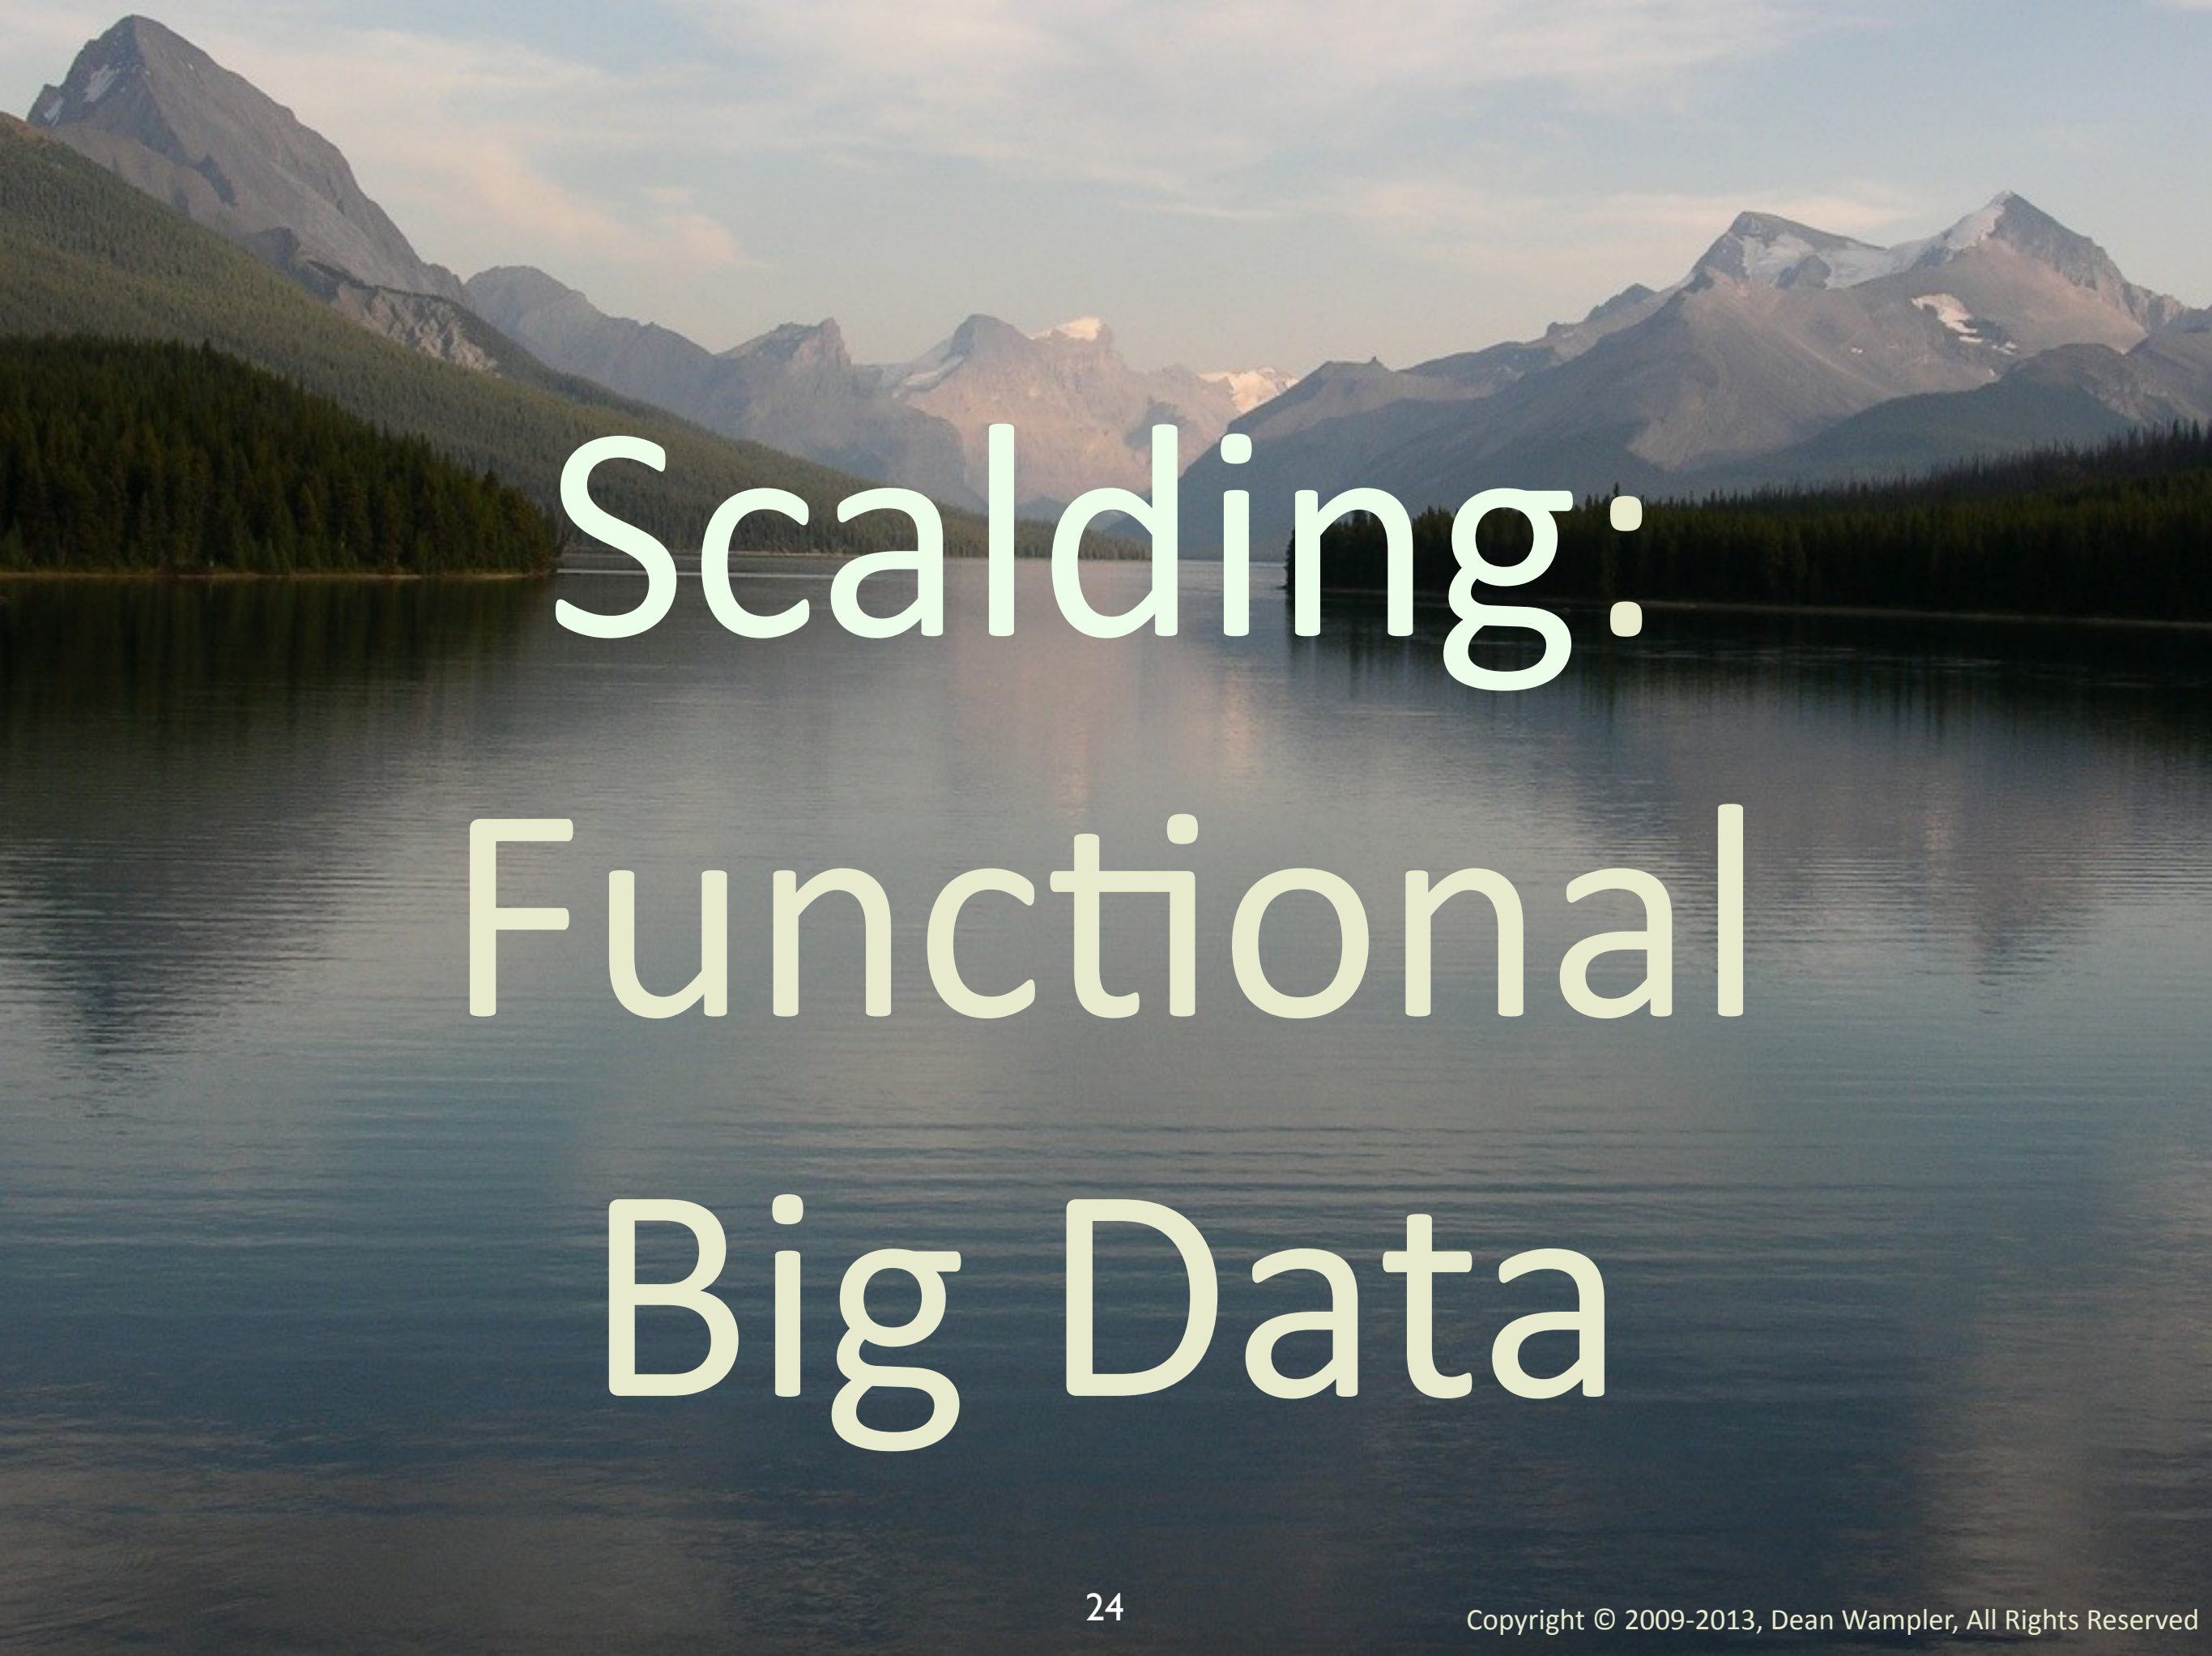
Variations on an Actor
theme.

drawing-actors.scala



Refine Features

- Add more geometric shapes
- Add other messages for new behaviors.
- Redraw all shapes drawn so far.
- Clear the “screen”.
- ... (see the README).



Scalding: Functional Big Data

Big Data

Data so big that
traditional solutions are
too slow, too small, or too
expensive to use.

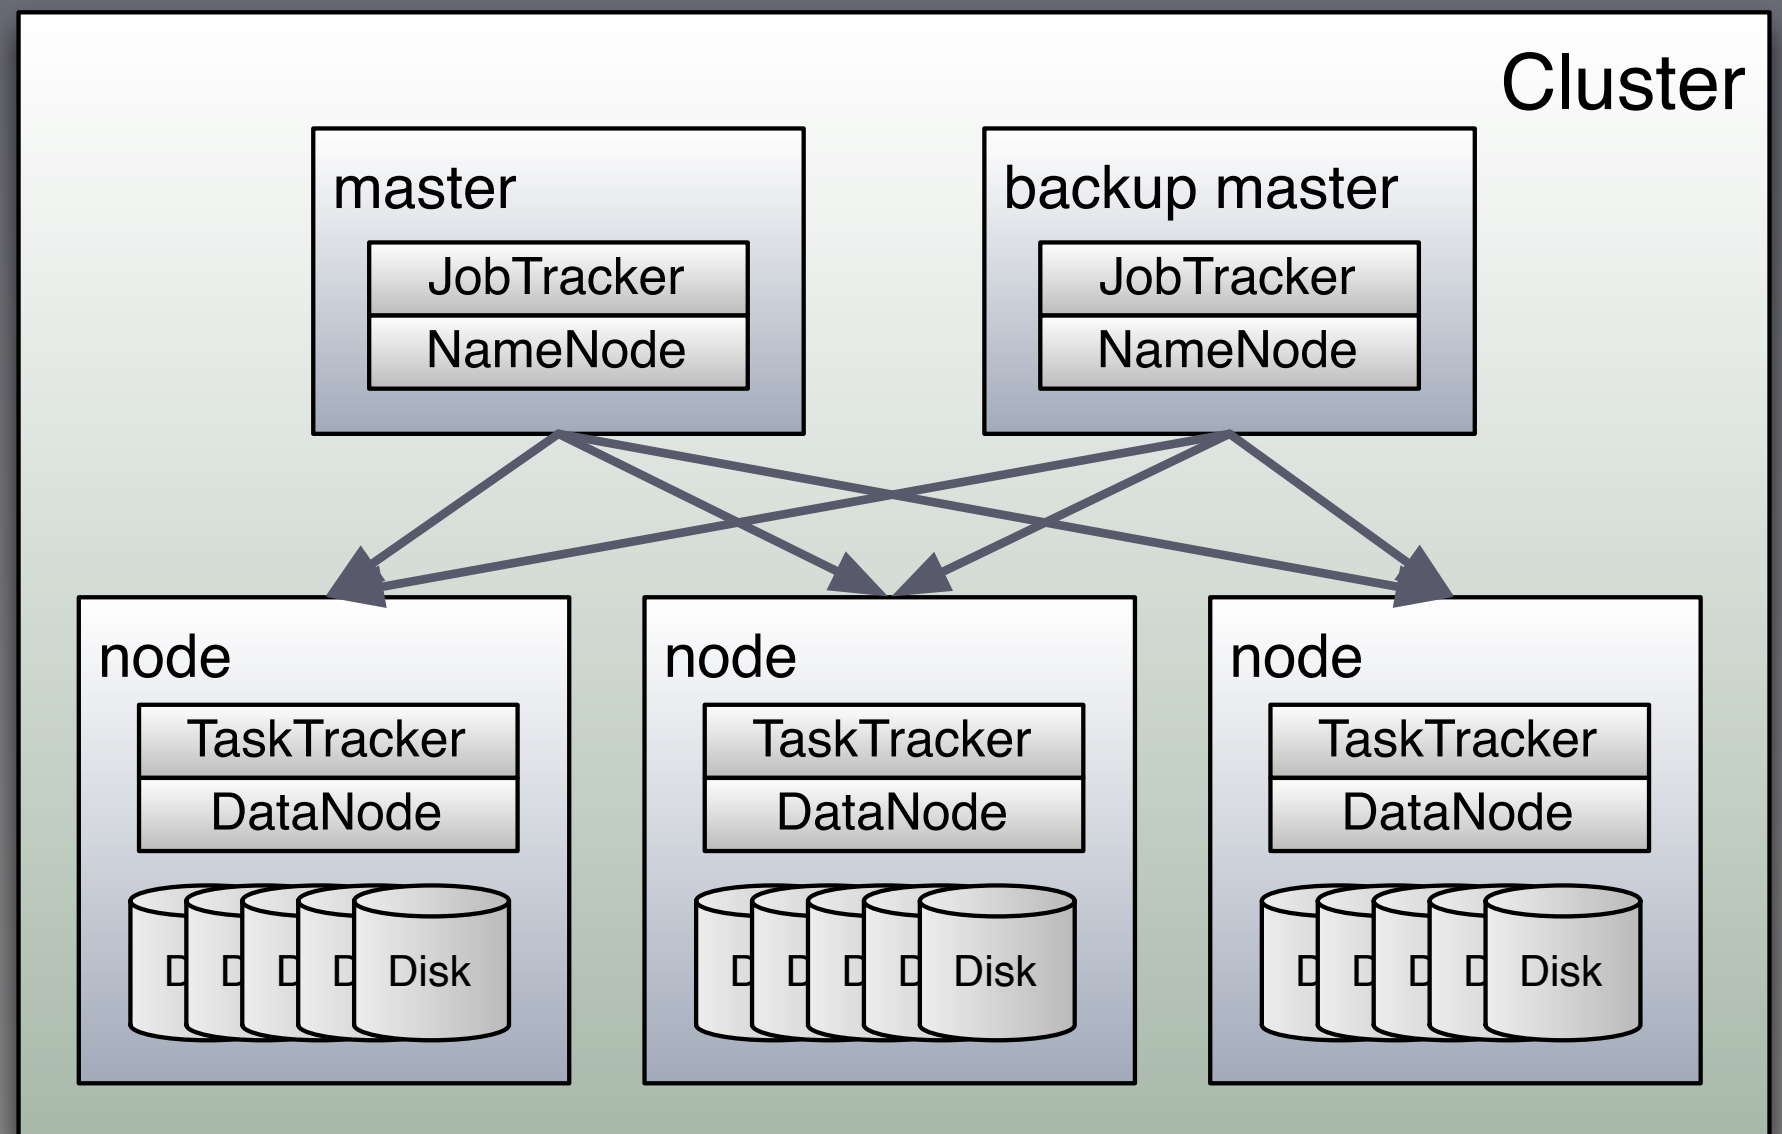


Hat tip: Bob Korbus

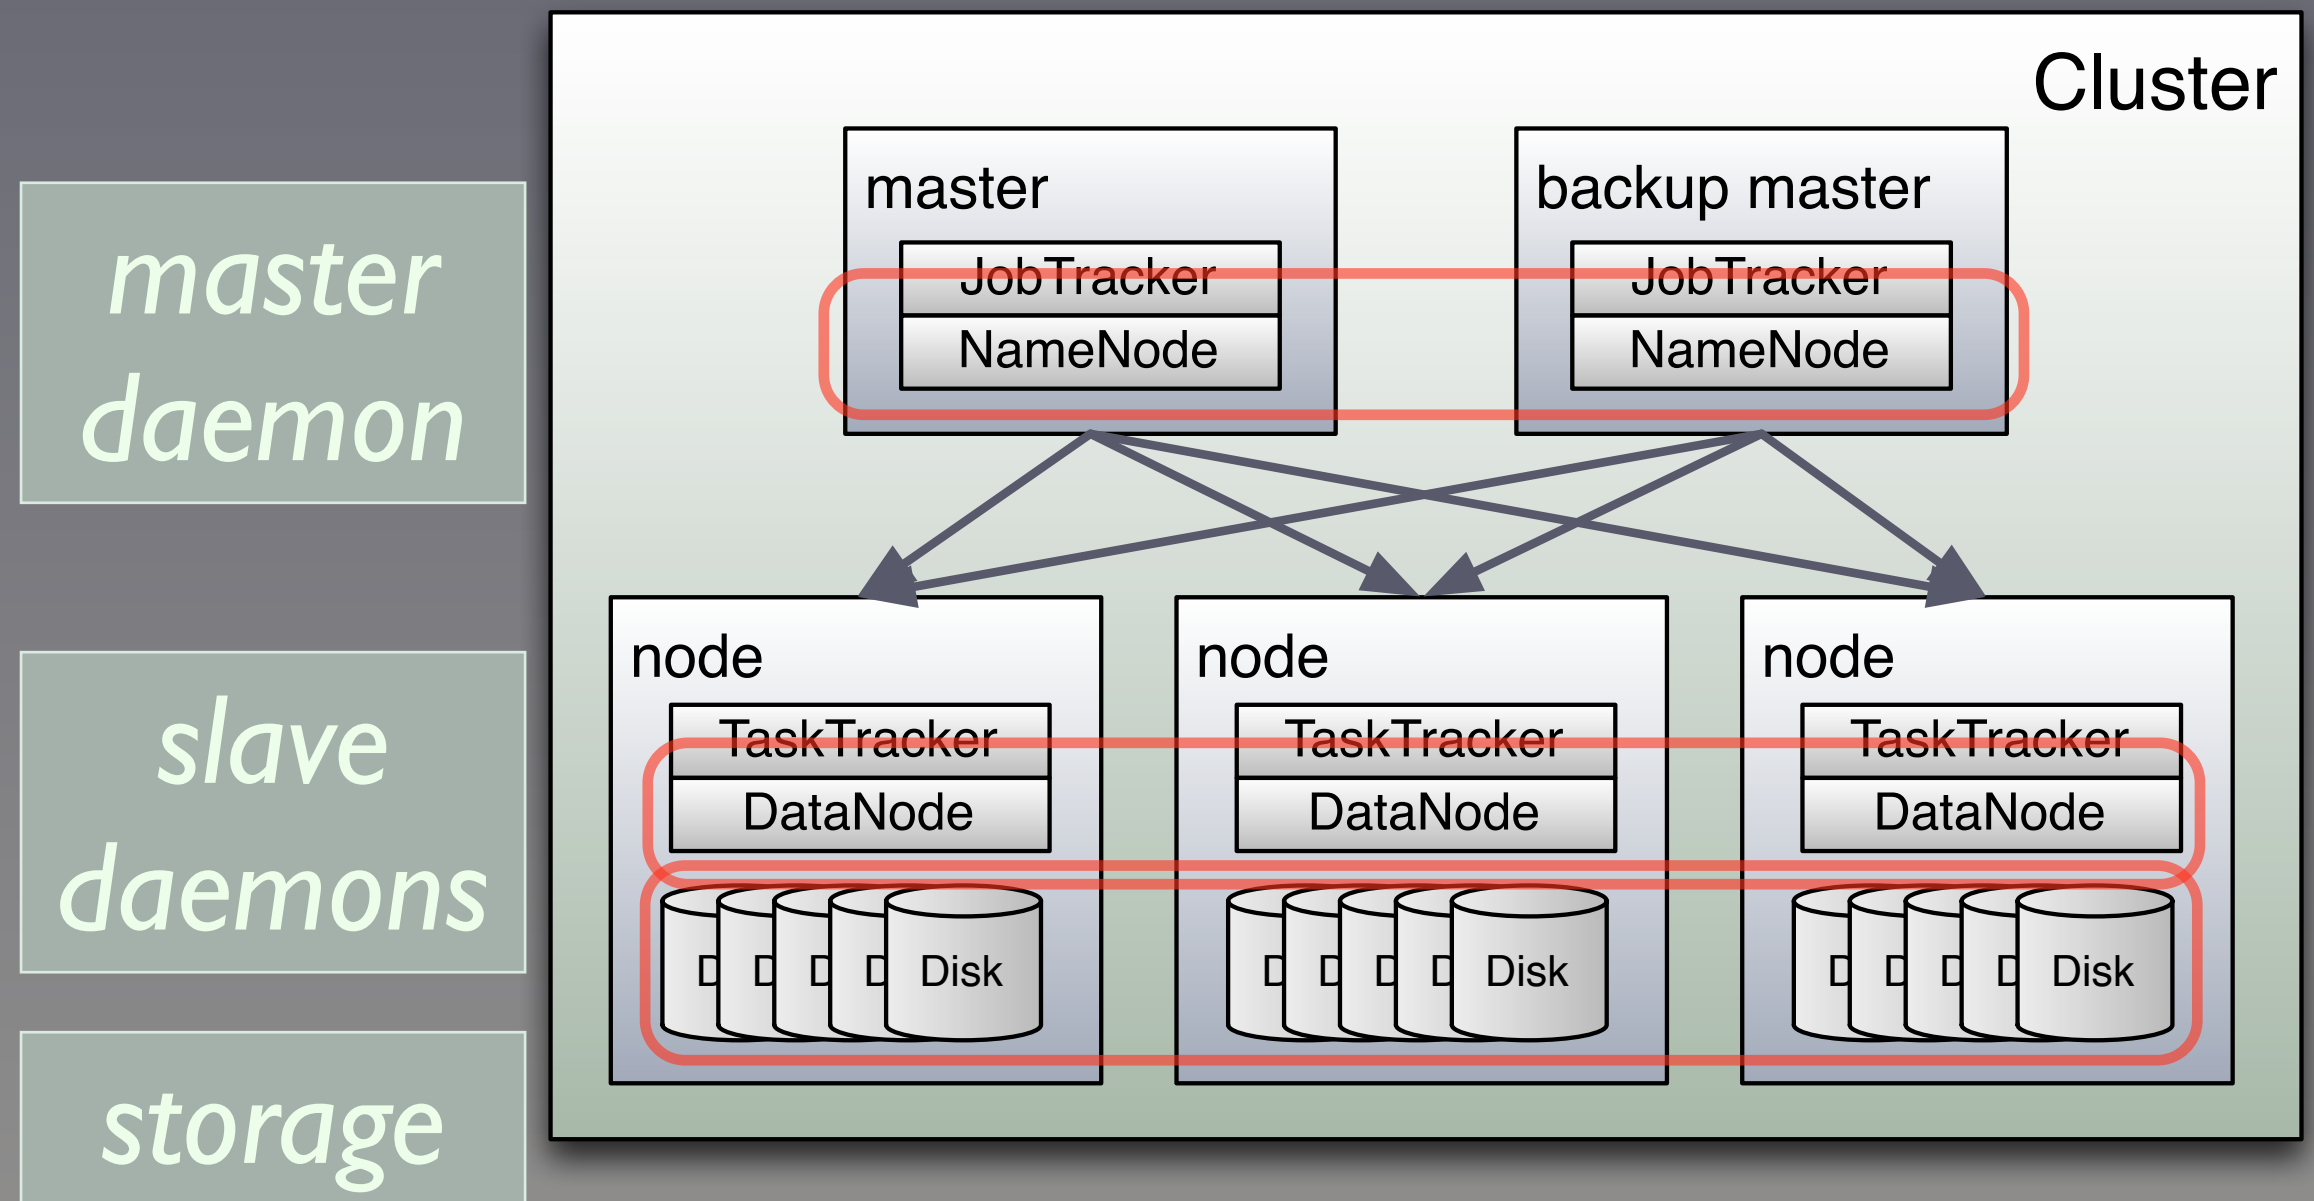
Hadoop Crash Course

- The most popular, general-purpose, distributed data framework.

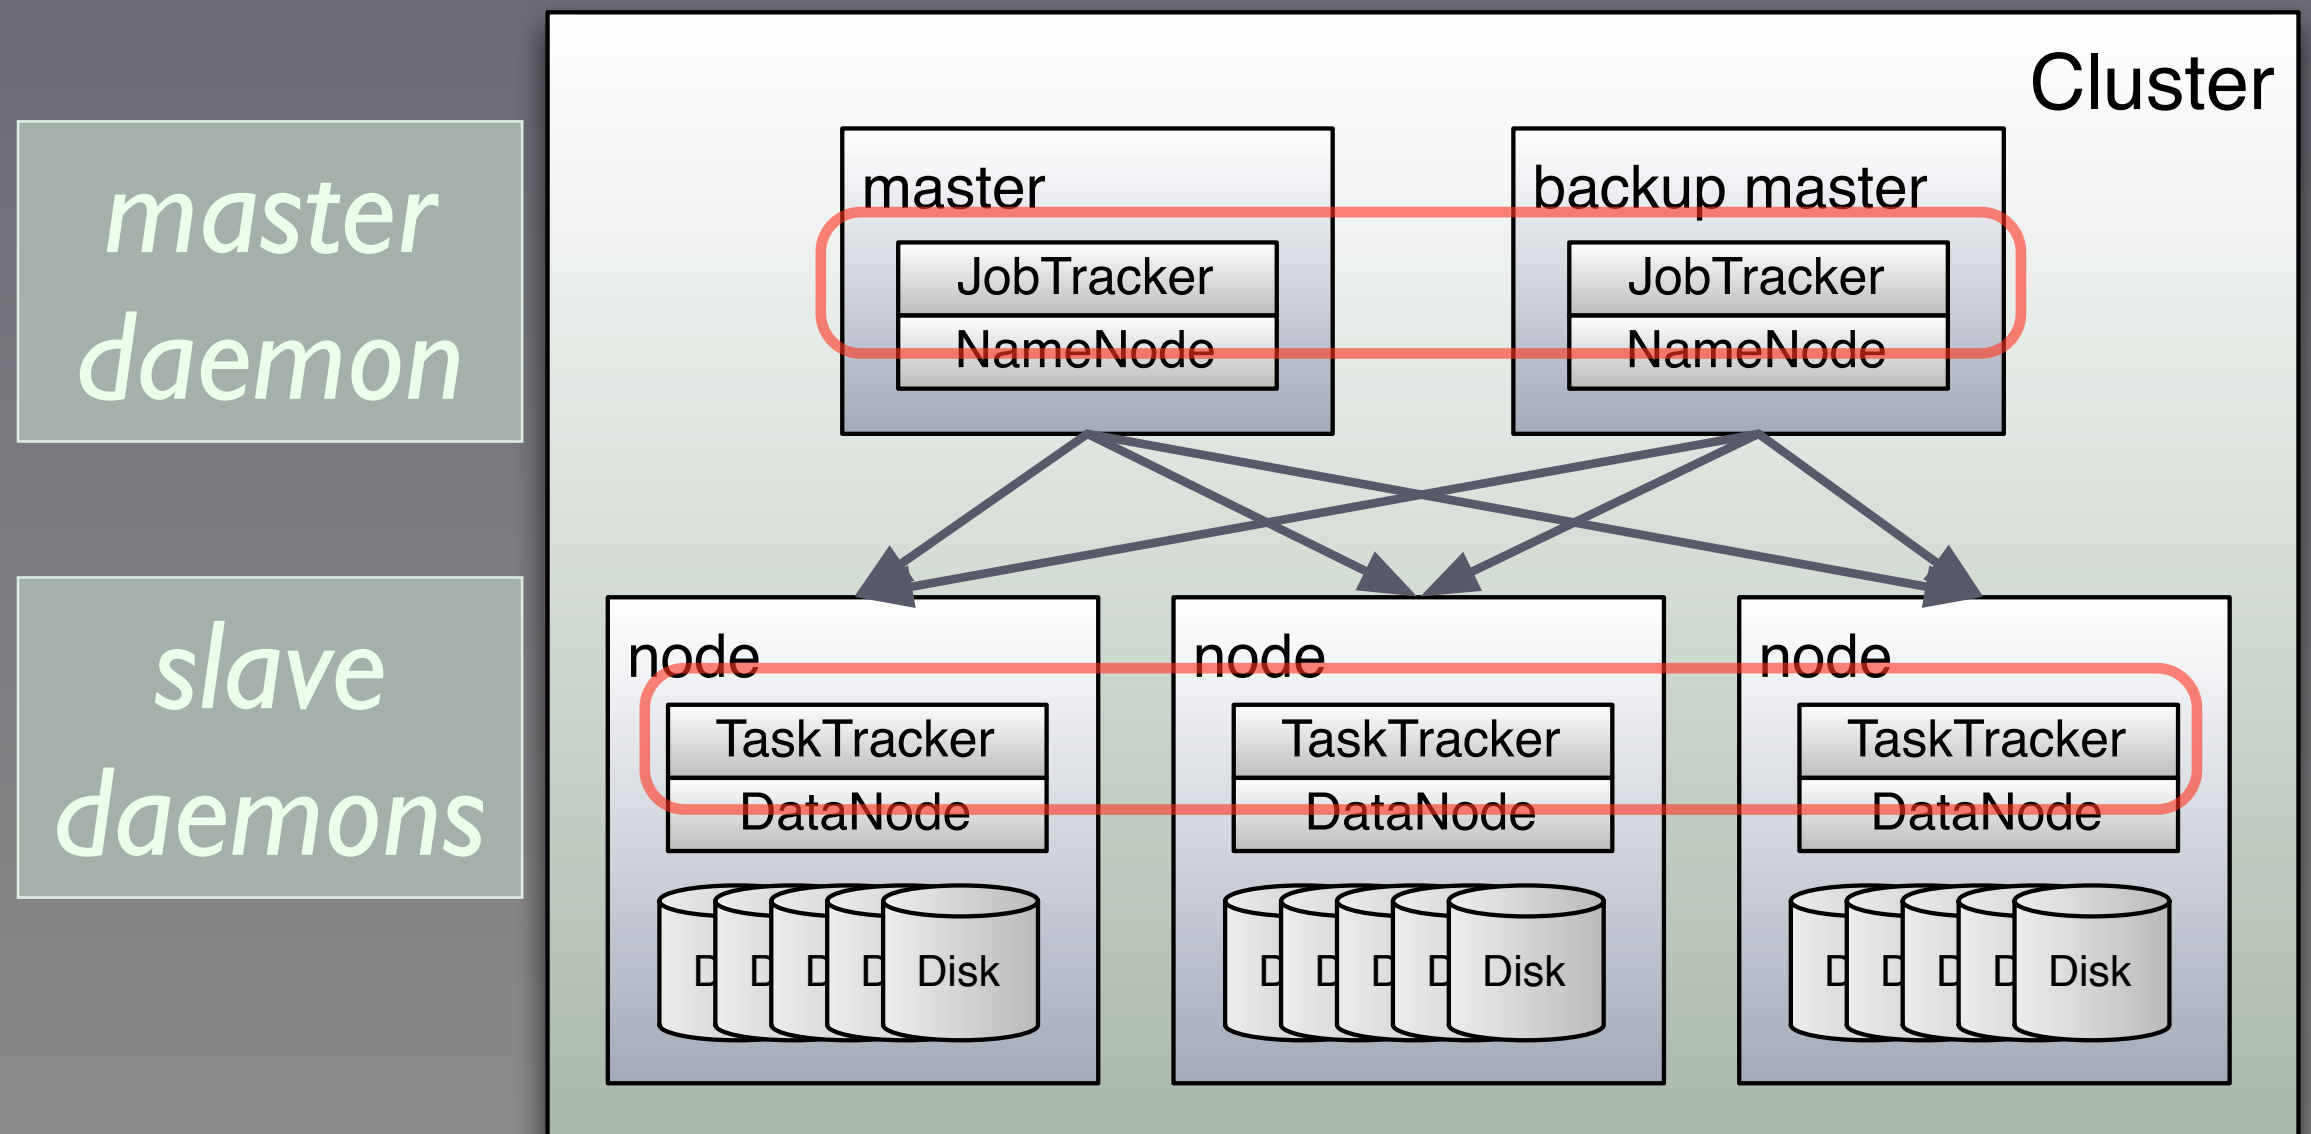
Hadoop Crash Course




Hadoop Distributed File System (HDFS)



Hadoop Jobs



What is MapReduce?



30

Copyright © 2009-2013, Dean Wampler, All Rights Reserved

Monday, April 22, 13

The famous Word Count algorithm. We have a corpus of documents on the left. We want to parse and count the words. By default, each document gets a Map task (JVM process) - the rounded bubbled, which we have to implement. Each line of the file is passed in with its position offset as the "key". We discard the key, tokenize the line, and emit new key-value pairs, each word and a count of 1, for every word found. (There are obvious optimizations to you could use...). Hadoop sorts the keys within each mapper task, then "shuffles" the key-value pairs over the cluster network to the appropriate reducer task (which we also must implement). We have set up the example so that all keys that start with 0 through 9 or a through l go to the 1st reducer, etc. Finally, each reducer is passed the keys, one at a time, with a collection of the counts for each key. The reducer sums the collections and writes the output, at least one file per reducer.

Hadoop Languages and Tools

- Hive: SQL queries.
- Pig: Dataflow language.
- Java API: “Assembly language”.
- Cascading: High-level Java API.
- Scalding: Scala Cascading API.

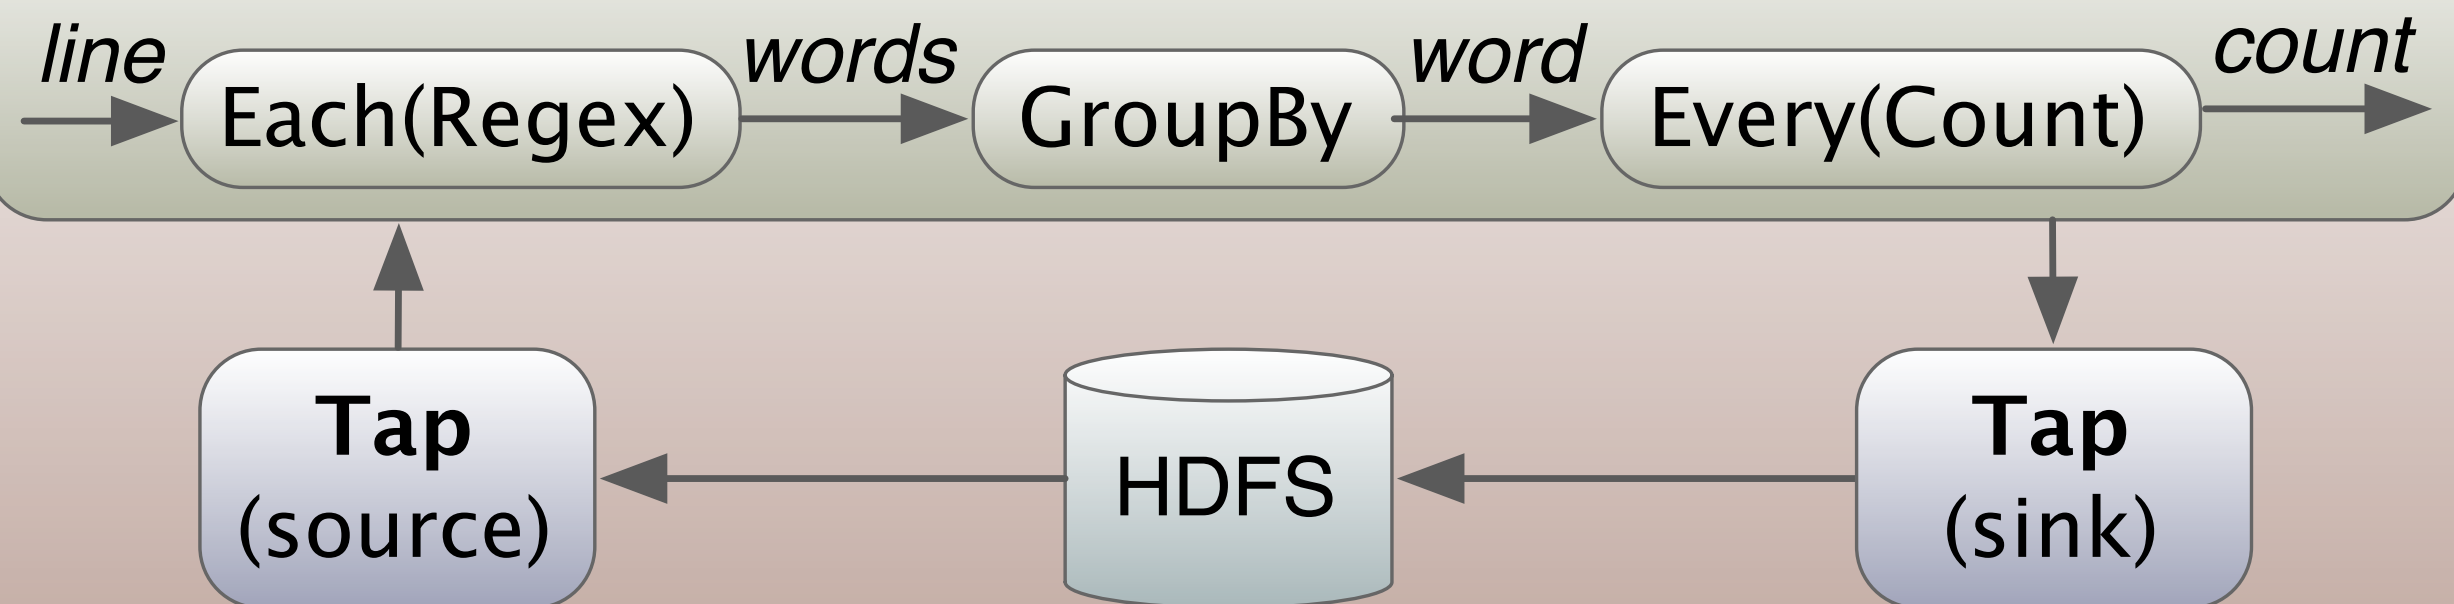
Cascading

Data flows consist of
source and sink Taps
connected by Pipes.

Cascading

Flow

Pipe ("word count assembly")



```

import org.cascading.*;
...
public class WordCount {
    public static void main(String[] args) {
        String inputPath = args[0];
        String outputPath = args[1];
        Properties properties = new Properties();
        FlowConnector.setApplicationJarClass( properties, Main.class );

        Scheme sourceScheme = new TextLine( new Fields( "line" ) );
        Scheme sinkScheme = new TextLine( new Fields( "word", "count" ) );
        Tap source = new Hfs( sourceScheme, inputPath );
        Tap sink = new Hfs( sinkScheme, outputPath, SinkMode.REPLACE );

        Pipe assembly = new Pipe( "wordcount" );

        String regex = "(?<!\pL)(?=\pL)[^ ]*(?<=\pL)(?!\\pL)";
        Function function = new RegexGenerator( new Fields( "word" ), regex );
        assembly = new Each( assembly, new Fields( "line" ), function );
        assembly = new GroupBy( assembly, new Fields( "word" ) );
        Aggregator count = new Count( new Fields( "count" ) );
        assembly = new Every( assembly, count );

        FlowConnector flowConnector = new FlowConnector( properties );
        Flow flow = flowConnector.connect( "word-count", source, sink, assembly);
        flow.complete();
    }
}

```

Here is the Cascading Java code. It's cleaner than the MapReduce API, because the code is more focused on the algorithm with less boilerplate, although it looks like it's not that much shorter. HOWEVER, this is all the code, where as previously I omitted the setup (main) code. See <http://docs.cascading.org/cascading/1.2/userguide/html/ch02.html> for details of the API features used here; we won't discuss them here, but just mention some highlights.

Note all the green types, but many represent “functions” to apply to the data, such as Count, Each, GroupBy, etc. I won't show the corresponding low-level Java API code, but Cascading does a good job emphasizing composition of behaviors and minimizing framework boilerplate.

Scalding

Scala API on top of
Cascading.

Developed at Twitter.

```

import com.twitter.scalding._

class WordCountJob(args: Args) extends
Job(args) {
  TextLine( args("input") )
    .read
    .flatMap('line -> 'word) {
      line: String =>
        line.trim.toLowerCase
          .split("\\W+")
    }
    .groupBy('word) {
      group => group.size('count) }
    }
    .write(Tsv(args("output")))
}

```

That's it!

Exercise

Scalding examples:
scalding-examples



Final Exercise

- Pick either the Akka or Scalding workshop/tutorial on GitHub and work through it.
- github.com/henrikengstrom/akka-meetup-sthlm
- github.com/deanwampler/scalding-workshop

A wide-angle photograph of a sandy beach. In the foreground, a large, dark, weathered piece of driftwood lies horizontally across the sand. The beach is vast and sandy, with some smaller pieces of driftwood scattered around. In the background, the ocean is visible with white-capped waves breaking. Two people riding horses are visible on the beach: one on the left side, further back, and another on the right side, closer to the water. The sky is a uniform, overcast grey.

Recap

Akka is a powerful toolkit
for distributed
applications.

Scalding is a great
example of why FP
is ideal for
data problems.

Thanks!

dean@concurrentthought.com
[@deanwampler](https://twitter.com/deanwampler)

polyglotprogramming.com/
programmingscala.com

