

@DEANWAMPLER

CHICAGO SPARK, MAY 18, 2016

SPARK 2.0

MEETUPS

- ▶ We need speakers and topics for future meetings!
- ▶ Thanks to Expedia for hosting...
 - ▶ ... and congratulations on your 20 year anniversary!

GOTO CHICAGO IS NEXT WEEK!!

► <http://gotocon.com/>

SPARK SUMMIT SF, JUNE 6-8

- ▶ spark-summit.org/2016/
- ▶ Discount codes:
 - ▶ Meetup16SF N% (?)
 - ▶ reynold16 20%

ABOUT YOU...

- ▶ Who's looking?
- ▶ Who's hiring?
- ▶ What topics should we cover soon?



SPARK 2.0.0

A NEW HOPE?

LIBERALLY BORROWED FROM...

- ▶ <http://go.databricks.com/apache-spark-2.0-presented-by-databricks-co-founder-reynold-xin>
- ▶ You should really view this webinar...

WHAT DOES A 2.0.0 RELEASE MEAN?

- ▶ Some new and restructured APIs.
- ▶ Some breaking API changes?
 - ▶ They try *very hard* to avoid changing user-visible APIs.
 - ▶ Some dependency changes, e.g., Guava.
 - ▶ Changes in experimental APIs (e.g., Datasets).

THREE MAJOR CHANGES

- ▶ Tungsten Phase 2
 - ▶ 5x-20x additional performance improvements.
- ▶ SQL 2003 and Unified DataFrames/Datasets API.
- ▶ Structured Streaming
 - ▶ Integration of DataFrames/Datasets and Streaming.



TUNGSTEN

PHASE 2

FOR AN ADDITIONAL 10X IMPROVEMENT...

- ▶ Removing hot spots only gains a few %.
- ▶ Ask instead, if we start from scratch, what's the fastest it could be?

```
select count(*) from sales  
where sku = 1234;
```

```
int count = 0;  
for (Record sale: sales) {  
    if (sale.sku == 1234)  
        count += 1;  
}
```

FOR AN ADDITIONAL 10X IMPROVEMENT...

- ▶ Instead, databases (and Spark) use the “volcano pattern”, where each filter is actually an iterators.
- ▶ Infrastructure adds overhead, easily 10x.


FOR AN ADDITIONAL 10X IMPROVEMENT...

- ▶ **After** compiling the query, generate custom code!
 - ▶ No virtual function calls.
 - ▶ Put data in CPU registers, when possible.
 - ▶ Unroll loops.
 - ▶ Exploit parallelism and pipelining.

RESULTS

Operation	Spark 1.6	Spark 2.0
filter	15ns	1.1ns
sum with/without group by	79/14ns	10.7/0.9ns
hash join	115ns	4.0ns
sort (8/16 bit entropy)	620/620ns	5.3/8.0ns
sort-merge join	750ns	700ns

PRELIM. TPC-DS BENCHMARK (LOWER IS BETTER)





EXAMPLE

- ▶ <http://bit.ly/1X8LKmH>
- ▶ (Example from Databricks Cloud)



DATASET, DATAFRAME

SQL

SQL 2003 COMPLIANCE

- ▶ Spark SQL now supports all 99 TPC-DS benchmark queries.
- ▶ New SQL parser (with better error messages).
- ▶ Also
 - ▶ Subqueries, correlated and uncorrelated.
 - ▶ Approximate aggregate statistics.

DATASETS VS. DATAFRAMES

- ▶ 2015: Datasets bring *field* type safety back to DataFrames.
 - ▶ Like static type safety provided by the RDD API:
 - ▶ `Dataset[T]` analogous to `RDD[T]`, where `T` is the record type.
 - ▶ In DataFrames, fields (columns) untyped => `Row` type.
- ▶ Now: `DataFrame = Dataset[Row]`

SPARKSESSION

- ▶ `SparkSession`: The new “`SparkContext`” for `DataFrame/Dataset`.
- ▶ Entry point for ingesting data (like `SQLContext` was).
- ▶ Metadata, configuration, and cluster resource management.

FUTURE?

- ▶ **RDD**: Will remain the low-level API.
- ▶ But **Dataset/DataFrame** will be the focus of optimizations, rich semantics.
- ▶ Higher-level libraries will be migrated to **Dataset/DataFrame**:
 - ▶ Structured streaming
 - ▶ ML pipeline replacing MLlib
 - ▶ GraphFrames

OTHER, MISCELLANEOUS API IMPROVEMENTS

- ▶ ML pipeline coverage in all languages (Java, Python, R, as well as Scala) nearly complete.
- ▶ Improved R support:
 - ▶ Parallelizable user-defined functions in R.
 - ▶ More Models!
 - ▶ Generalized Linear Models (GLMs), Naïve Bayes, Survival Regression, K-Means.

EXAMPLES

- ▶ <http://bit.ly/1SMPEzQ>
- ▶ <http://bit.ly/1OeqdSn>
- ▶ (Examples from Databricks Cloud)



STRUCTURED

STREAMING

WHY STREAMING?

- ▶ Spark started as a batch mode system, but...
 - ▶ ... the mini batch model lets Spark capture data in small time windows and run batch jobs over it.
- ▶ This also lets you use business logic in both batch and streaming contexts, which is great.
- ▶ If you are not latency sensitive, mini batch lets you do fancy things:
 - ▶ “Online” training of ML models, track moving state, run SQL queries, ...

STREAMING IS HARD

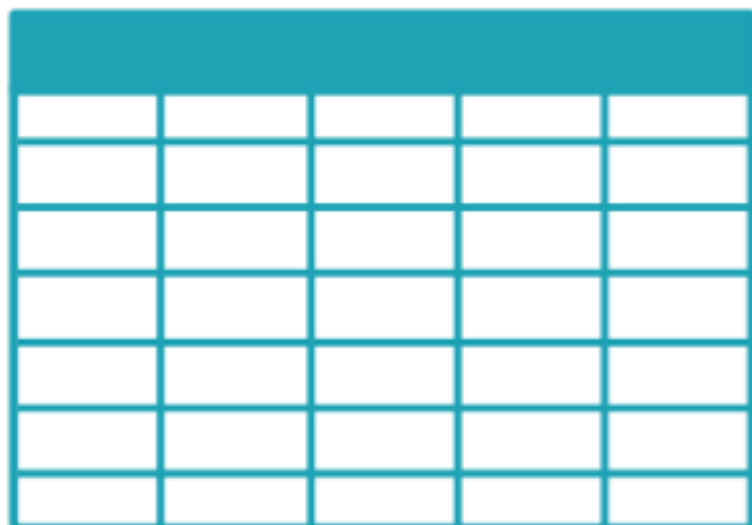
- ▶ I'm counting sales/hour, but some messages get delayed for several hours. Now what?
- ▶ Event time vs. processing time.
- ▶ What does group by or join mean in a streaming context?
- ▶ If I'm maintaining stream state and it crashes, how do I recover?

STREAMING IS HARD

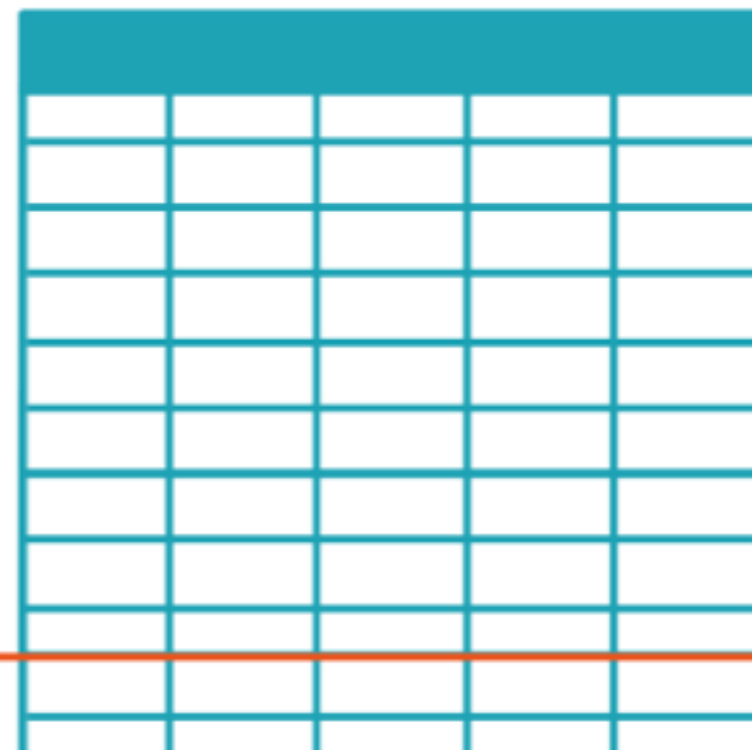
- ▶ Semantics are nontrivial:
 - ▶ <https://www.oreilly.com/ideas/the-world-beyond-batch-streaming-101>
 - ▶ <https://www.oreilly.com/ideas/the-world-beyond-batch-streaming-102>
 - ▶ Excellent [Kafka Summit talk](#) by Frances Perry and Tyler Akidau
 - ▶ ... and others.

GOAL: ELIMINATE THE NEED TO REASON ABOUT STREAMING

Spark 1.3
Static DataFrames



Spark 2.0
Infinite DataFrames





Single API !

CLASSIC BATCH JOB

```
val logs = ctx.read.format("json").  
  open("s3://mybucket/logs")  
  
logs.groupBy(logs.user_id).  
  agg(sum(logs.time)).  
  write.format("jdbc").  
  save("jdbc:mysql://...")
```

... CONVERTED TO A CONTINUOUS AGGREGATION

```
val logs = ctx.read.format("json").  
  stream("s3://mybucket/logs")
```

```
logs.groupBy(logs.user_id).  
  agg(sum(logs.time)).  
  write.format("jdbc").  
  startStream("jdbc:mysql//...")
```

- ▶ `open` changed to `stream`.
- ▶ `save` changed to `startStream`.

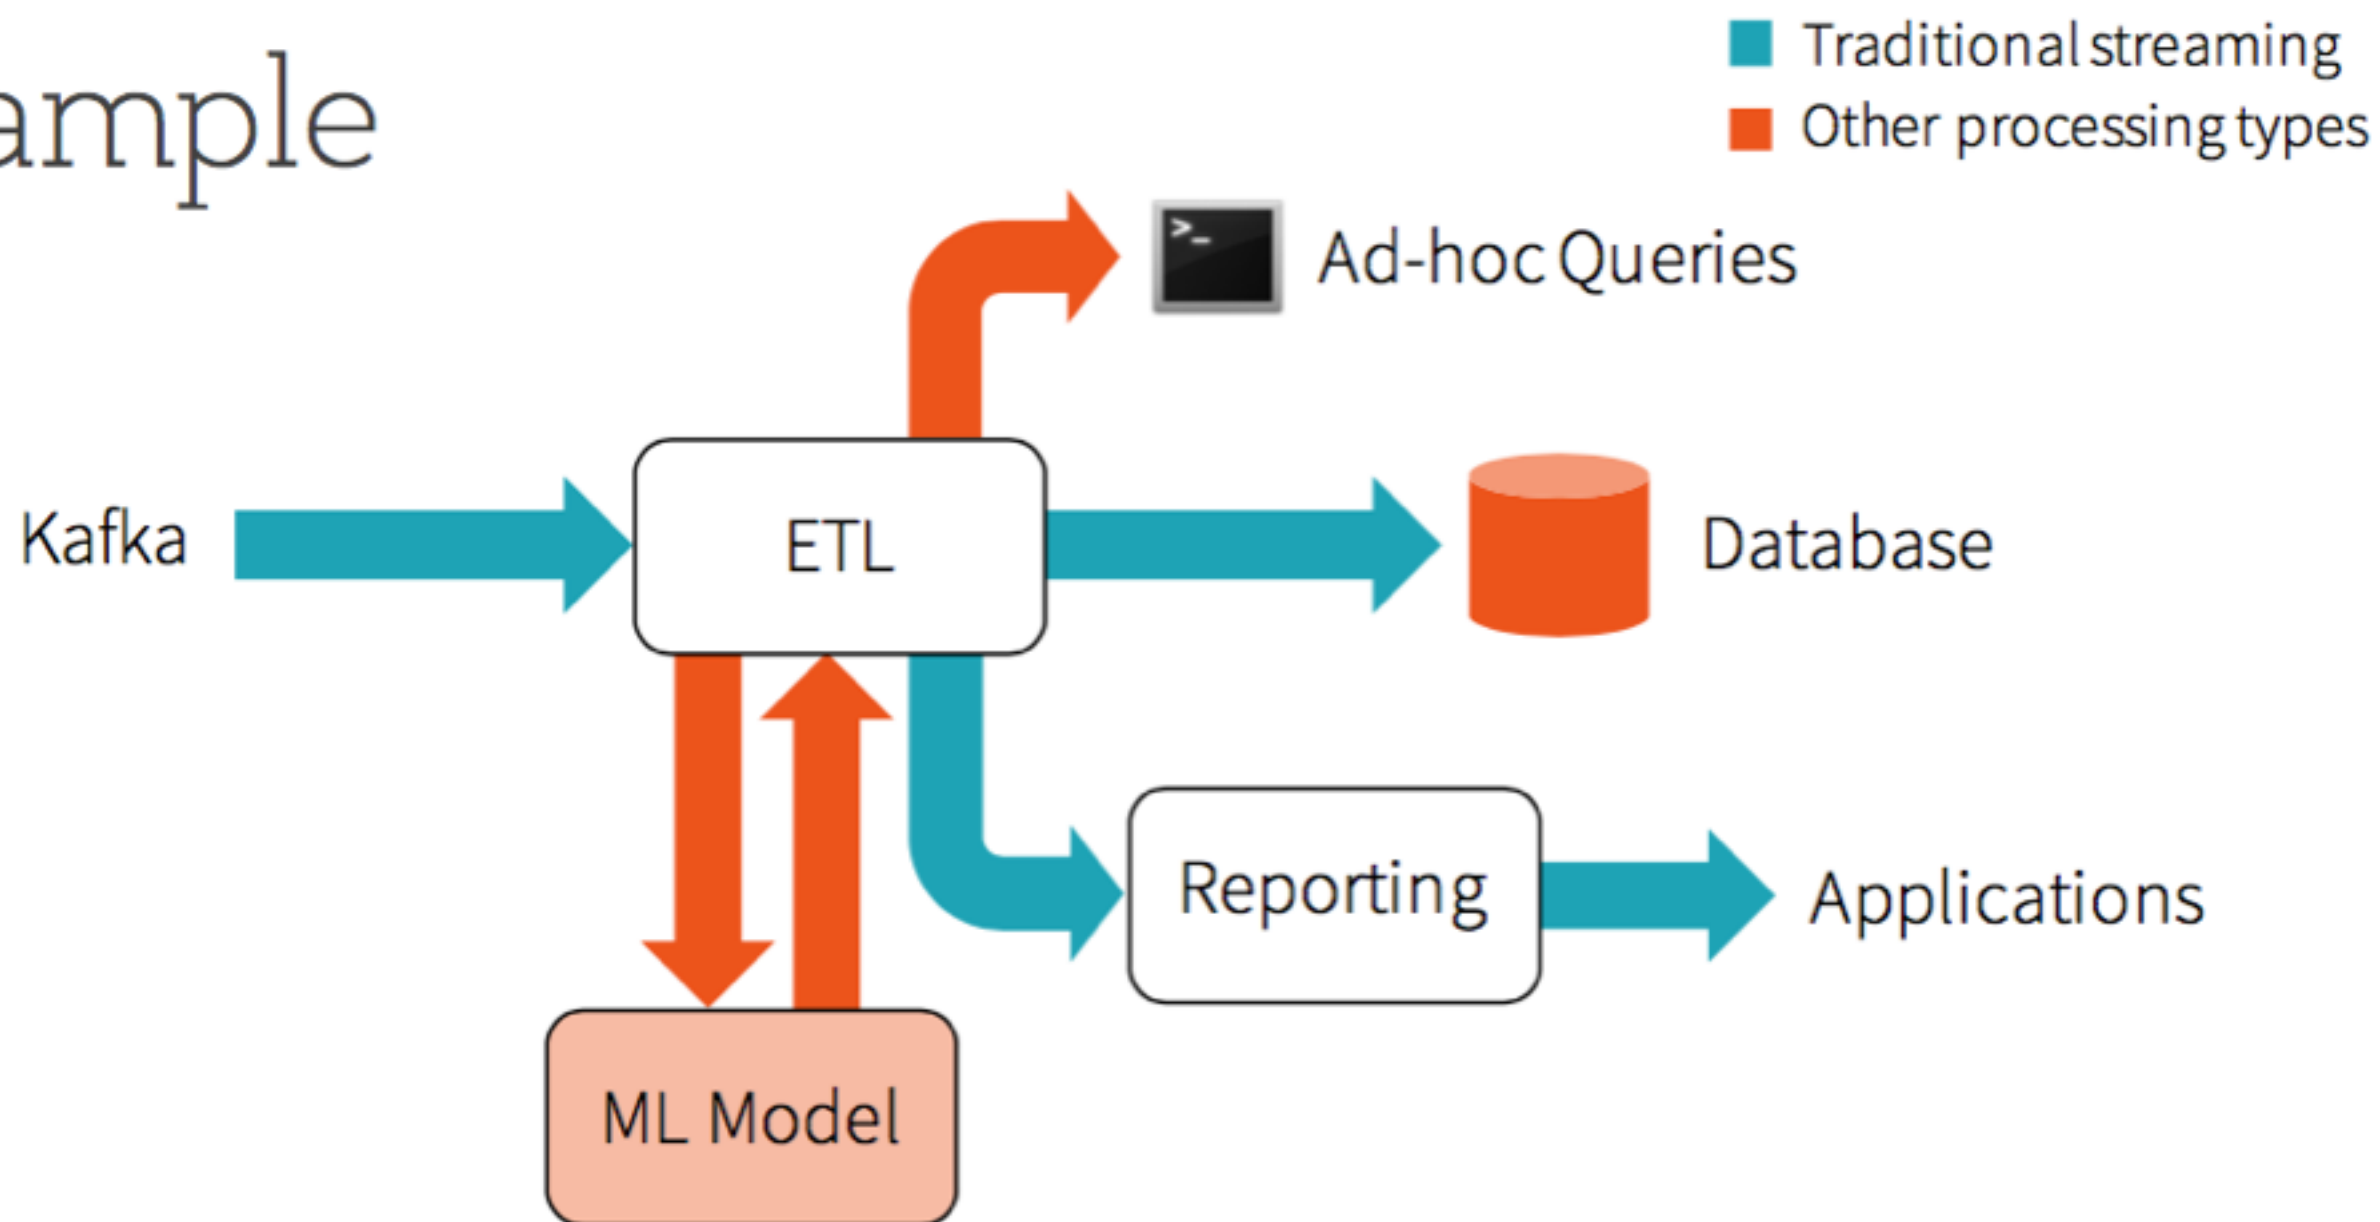
SO, STRUCTURED STREAMING IS...

- ▶ High-level streaming API built on SparkSQL engine, rather than RDDs.
- ▶ Supports:
 - ▶ Event time vs. processing time.
 - ▶ Windowing.
 - ▶ Sessions.
 - ▶ Misc. sources and sinks.

SO, STRUCTURED STREAMING IS...

- ▶ Queries:
 - ▶ The usual SQL-like aggregations, etc.
 - ▶ Query the stream state using JDBC.
 - ▶ Change queries at runtime.
 - ▶ Build and apply machine learning models.
- ▶ Build “continuous applications”.

Example



Goal: end-to-end continuous applications

FUTURE?

- ▶ There is talk of rewriting Spark Streaming to be a “true” streaming engine.
 - ▶ Low latency.
 - ▶ Full support for common streaming semantics
 - ▶ (i.e., as discussed in the Tyler Akidau blogs.)
- ▶ Needs to stay competitive with [Apache Beam](#), [Flink](#), and [Gearpump](#).



FOR MORE

INFORMATION

RESOURCES

- ▶ Structured Streaming Strata Talk
 - ▶ <https://www.oreilly.com/learning/apache-spark-2-0--introduction-to-structured-streaming>
- ▶ 2.0 Preview:
 - ▶ code: <http://home.apache.org/~pwendell/spark-releases/spark-2.0.0-preview-bin/>
 - ▶ docs: <http://home.apache.org/~pwendell/spark-releases/spark-2.0.0-preview-docs/>
 - ▶ Databricks Cloud Preview.

RESOURCES

- ▶ lightbend.com/fast-data
- ▶ dean.wampler@lightbend.com
- ▶ [@deanwampler](#)
- ▶ polyglotprogrammin.com/talks
- ▶ Questions?